

# CHPC Student Cluster Competition 2022

## WRF

The **W**eather **R**esearch and **F**orecasting model is a state-of-the-art mesoscale **N**umerical **W**eather **P**rediction (**NW**P) system designed for both atmospheric research and operational forecasting applications. It features two dynamical cores, a data assimilation system, and a software architecture supporting parallel computation and system extensibility. The model serves a wide range of meteorological applications across scales from tens of meters to thousands of kilometers.

Obtain copies of WRF and its Pre-Processing System source codes from GitHub <https://github.com/wrf-model/WRF> and <https://github.com/wrf-model/WPS> (the benchmarks have been tested against **WRF 4.2**). Alternatively copies of the **WRF 4.0** tarballs have also been included in your competition folder.

Configuration, building and installation instructions can be found at <https://www2.mmm.ucar.edu/wrf/OnLineTutorial/>, which detail the following dependencies to be installed or made available in your environment:

- Your choice of **C** and **FORTRAN COMPILERS** and implementation of **MPI**
- **HDF5** library required for **NETCDF4** compression support. (Tested with version 1.10.5)
- Optionally you can make use of **PARALLEL HDF5**, which will process file I/O in parallel, as opposed to serially
- **NETCDF** library is used for large network file I/O with gridded model data. (Tested with version 4.6.3)
- Optionally you can make use of **PARALLEL HDF5** and **NETCDF** libraries, which will process network file I/O in parallel, as opposed to serially
- If these are not already installed on your system, you may additionally require **JASPER**, **LIBPNG** and **ZLIB**.

Take care to ensure that you are working within a consistent environment. Failure to manage your environment, libraries and dependencies properly will impede your team from being able to successfully build and compile the application. You can compile for **BASIC NESTING**.

During the compilation process you will be presented with a choice of options to select for a column (type of compiler and architecture) and a row (type of parallel build):

- The first column is for **SERIAL** (single processor) builds.
- The second column is **OPENMP** (threaded, shared-memory), builds with up to 40 maximum processes.
- The third and fourth columns are for (**MPI** only) and (**OPENMP + MPI**)

In an analogous way to the optimization experiment from the **SWIFT** benchmark, where you optimized for the most efficient communication pattern between **MPI** processes, with each containing multiple **POSIX** (**P****T****H****R****E****A****D****S**) utilizing shared memory. In this experiment however, you will be required to find an optimal

mix of **MPI** process and **OPENMP** threads, (or your configuration may perform best without **OPENMP** threads).

Clean build the executable on multiple processors, and for debugging purposes you can usefully redirect and store standard out and standard err to a log file.

```
$ ./compile - j N em_real >& build_wrf.log
$ ./configure
```

Once **WRF** has finished building, you must compile and configure **WPS**, and from within this subdirectory, you may need to create the following symbolic link:

```
$ ln -sf ungrib/Variable_Tables/Vtable.GFS Vtable
```

## Benchmark 1: Single Domain Case

Extract the **GEOG** and **DATA** into your **WRF/GEOG** and **WRF/DATA** folders, respectively, then copy **namelists.wps** to **WRF/WPS**. From your **WPS** folder, link the weather data files and build the geographical land usage interpolation and metrological data:

```
$ ./link_grib.csh ../DATA/*
$ ./ungrib.exe >& ungrib.log
$ ./geogrid.exe >& geogrid.log
$ ./metgrid.exe >& metgrid.log
```

Before you can run the benchmark, you must copy the benchmark input files **namelist.input** into the **run/ benchmark** subdirectory, create symbolic links to the metrological input data and execute the last interpolation step. From the **run/ benchmark** subdirectory:

```
$ ln -sf PATH/TO/met_em.d01*
$ ./real.exe >& real.log
```

This will generate the boundary condition file for the (outer) domain `wrfbdy_d01`, and the initial conditions for the domain `wrfinput_d01`, (there will be an additional input file in the case for nested domains). There will be a `wrf.exe` binary and a symbolic link to them under the `main/` and `run/` folders, respectively.

You must determine the optimum ratio of **OPENMP THREADS** to **MPI RANKS**, and in this regard useful environment variables to manipulate include `OMP_STACKSIZE` and `OMP_NUM_THREADS`. Configure an appropriate SLURM batch script and reroute the output to an appropriately named logfile.

```
time mpirun -genv OMP_NUM_THREADS <OpenMP Threads> \
            -np <Num Nodes><MPI Ranks> \
            2>&1 | tee singleDomain.log
```

You will notice `rsl.error.xxxx` and `rsl.out.xxxx` files, corresponding to each of the `xxxx` MPI ranks. It may be useful to monitor or `tail` the output of the `.0000` logs. Submit the `rsl.error`, `rsl.out`, `singleDomain.log` logs and the SLURM batch configuration file.

## Benchmark 2: Nested Two-Domain Case

Repeat the above experiments, to optimize the ratio of OpenMP Threads to MPI Ranks for the nested Two-Domain Case. Please take note that you do not need to recompile the application, unless you are optimizing your application binaries through:

- Testing various implementations of Compiler and MPI, and/or
- Dependency and library versions, and/or
- Compiler optimization flags, and/or
- Any other means to change the application binary,